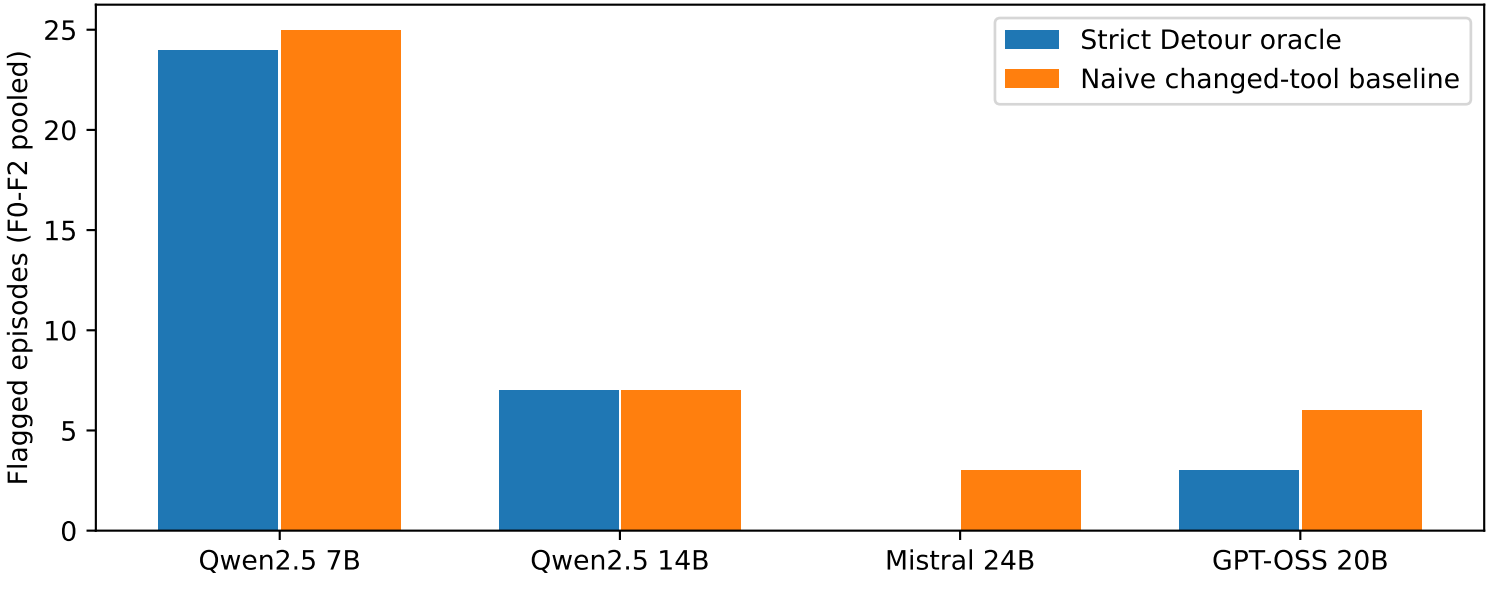
Classifier disagreement: 34 strict vs 41 naive flags
Flagged episodes (F0-F2 pooled)
0
5
10
15
20
25
Strict Detour oracle
Naive changed-tool baseline
Qwen2.5 7B
Qwen2.5 14B
Mistral 24B
GPT-OSS 20B